

linear-regression-lab

January 23, 2023

1 Linear Regression Lab

1.1 Load Data

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

df = pd.read_csv("housing.csv")
#df['price'] /= 1000
#df['sqft'] /= 1000
df
```

```
[ ]: # Set up our training data
x_train = df['sqft'].to_numpy()
y_train = df['price'].to_numpy()

# To use only a subset of the training data, use a slice of the arrays above.

print("Training data for x: ", x_train)
print("Training data for y: ", y_train)
```

```
[ ]: # Make a nice plot

plt.scatter(x_train, y_train)
plt.show()
```

```
[ ]: # Make a nice plot with a line based on the equation  $y = wx + b$ 

plt.scatter(x_train, y_train)

# Set up dummy values for  $w$  and  $b$ 
w = 0
b = 0

# Generate points on our line  $y = wx + b$  to graph
line_x_points = np.linspace(min(x_train), max(x_train), 100)
line_y_points = [w * x + b for x in line_x_points]
```

```
plt.plot(line_x_points, line_y_points, c='red')
plt.show()
```

```
[ ]: # Let's turn it into a function:
```

```
def make_plot(x_train, y_train, w, b):
    plt.scatter(x_train, y_train)

    line_x_points = np.linspace(min(x_train), max(x_train), 100)
    line_y_points = [w * x + b for x in line_x_points]

    plt.plot(line_x_points, line_y_points, c='red')
    plt.show()

make_plot(x_train, y_train, 0, 0)
```

```
[ ]: # Let's define a function to predict  $y = f(x) = wx + b$ 
```

```
def make_prediction(x, w, b):
    return w * x + b
```

```
[ ]: # Let's define a function to compute the cost  $J(w, b)$  on a set of data:
```

```
def compute_cost(x_data, y_data, w, b):
    """
    x_data and y_data are lists
    w and b are scalars
    return: gradient of w, gradient of b
    """
    m = len(x_data) # number of data points
    cost = 0

    for i in range(m):
        y_hat = make_prediction(x_data[i], w, b)
        y = y_data[i]
        cost += (y_hat - y)**2
    total_cost = (1 / (2 * m)) * cost

    return total_cost
```

```
[ ]: # Output total cost for our data set:
```

```
w = .3
b = 500

print(compute_cost(x_train, y_train, w, b))
```

```
[ ]: # Let's write a function to compute the gradient:
```

```
def compute_gradient(x_data, y_data, w, b):  
    """  
    x_data and y_data are lists  
    w and b are scalars  
    return: gradient of w, gradient of b  
    """  
    m = len(x_data) # number of data points  
    gradient_w = 0  
    gradient_b = 0  
  
    for i in range(m):  
        y_hat = make_prediction(x_data[i], w, b)  
        y = y_data[i]  
  
        gradient_w += (y_hat - y) * x_data[i]  
        gradient_b += (y_hat - y)  
  
    gradient_w /= m  
    gradient_b /= m  
  
    return gradient_w, gradient_b
```

```
[ ]: # Let's write code to run gradient descent:
```

```
w = 0  
b = 0  
ALPHA = .0000001  
  
J_sequence = []  
  
print("x train is", x_train)  
print("y train is", y_train)  
  
for ctr in range(0, 100):  
    print(ctr)  
    print(w, b)  
    print("Cost is", compute_cost(x_train, y_train, w, b))  
    (gradient_w, gradient_b) = compute_gradient(x_train, y_train, w, b)  
    print("Gradients", gradient_w, gradient_b)  
    w -= ALPHA * gradient_w  
    b -= ALPHA * gradient_b  
    J_sequence.append(compute_cost(x_train, y_train, w, b))  
  
print("Final w and b:", w, b)
```

```
[ ]: # Let's plot the cost as a function of number of iterations of the
      # gradient descent algorithm.

      plt.scatter(range(0, len(J_sequence)), J_sequence)
      plt.show()

[ ]: # Produce final plot of our data and our line:

      make_plot(x_train, y_train, w, b)

[ ]: # Make a new prediction:

      sqft = 2000
      prediction = make_prediction(sqft, w, b)

      print("A house with", sqft, "square feet would sell for around", prediction,
            ↪ "dollars.")

[ ]:
```